

Reinforcement Learning: Homework #1

Due on March 22, 2020 at 11:59pm

Professor Ziyu Shao

Tianyuan Wu
63305667

Problem 1

Sampling from probability distributions. Show histograms and compare them to corresponding PDF.

Solution

- 1) Sampling from the Logistic distribution by using Unif(0,1).

The CDF of Logistic distribution is: $F(x) = \frac{e^x}{1+e^x}$. So we can use inverse transform technique and find the inversion of it, which is $F^{-1}(u) = \log \frac{u}{1-u}$. So for each sample, we generate a uniform distributed r.v. and return the inversion of it. The python code is shown below.

```
def logistic():
    u = uniform(0, 1)
    return log(u/(1-u))
```

Also, we can find the PDF of the logistic distribution is $f(x) = \frac{e^{-x}}{(1+e^{-x})^2}$. Based on these analysis, I generate 200,000 sample and draw the histogram and PDF shown in Fig.1.




Figure 1: Logistic distribution via inverse transform

- 2) Sampling from the Rayleigh distribution by using Unif (0,1).

The CDF of Logistic distribution is: $F(x) = 1 - e^{-\frac{x^2}{2}}$. so we can also use the inverse transform to find $F^{-1}(u) = \sqrt{-2 \log(1-u)}$

```
def rayleigh():
    u = uniform(0, 1)
    return sqrt(-2*log(1-u))
```

Also, we can find the PDF of the logistic distribution is $f(x) = xe^{-\frac{x^2}{2}}$. Based on these analysis, I generate 200,000 sample and draw the histogram and PDF shown in Fig.2.

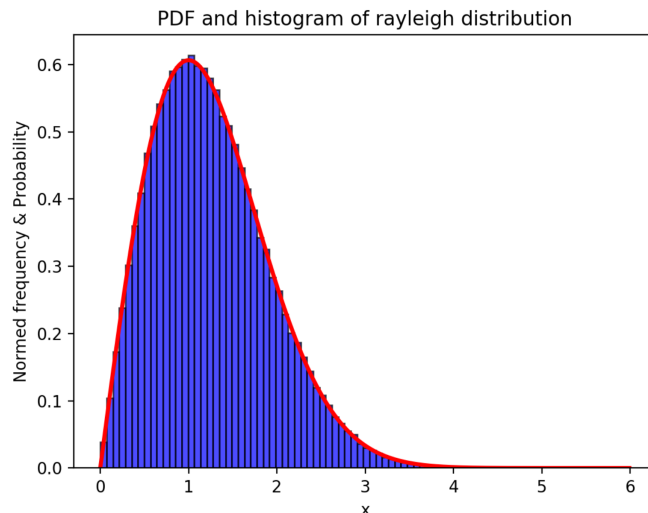


Figure 2: Rayleigh distribution via inverse transform

3) Sampling from the standard Normal distribution with both the Box-Muller method and the Acceptance-Rejection method. Compare the pros and cons of both methods.

1. Box-Muller method

First, generate two r.v. $u_1 \sim U(0, 1)$ and $u_2 \sim U(0, 1)$.

Then, by the Box-Muller method, $n_1 = \sqrt{-2 \log(u_1)} \cos(2\pi u_2)$ and $n_2 = \sqrt{-2 \log(u_1)} \sin(2\pi u_2)$ are 2 independent r.v. which both obey normal distribution.

```
def normal():
    u1 = uniform(0,1)
    u2 = uniform(0,1)
    return sqrt(-2*log(u1))*cos(2*pi*u2)
```

The result is shown in Fig.3




Figure 3: Normal distribution via Box-Muller method

2. Acceptance-Rejection method

First, we can generate a standard exponential distribution $q(x) \sim \text{Expo}(1)$ by inverse transform method.

Second, let $Z \sim \frac{2}{\sqrt{2\pi}} e^{-\frac{x^2}{2}} = p(x)$ which is half of the normal distribution.

Then, find $c = \sup_{\epsilon} \frac{q(\epsilon)}{p(\epsilon)}$, which equals to $\sqrt{\frac{2e}{\pi}}$

Finally, run the Acceptance-Rejection method, generate y from $q(x)$, and u from $U(0, 1)$. Accepts when $u < \frac{p(y)}{cq(y)}$. The python code is:

```
def normal_acc_rej():
    def expo_1():
        u = uniform(0, 1)
        return -log(u)
    y = expo_1()
    u = uniform(0, 1)
    if u < exp(-0.5*(y-1)**2):
        return y if random() < 0.5 else -y
    return normal_acc_rej()
```

The result is shown in Fig.4

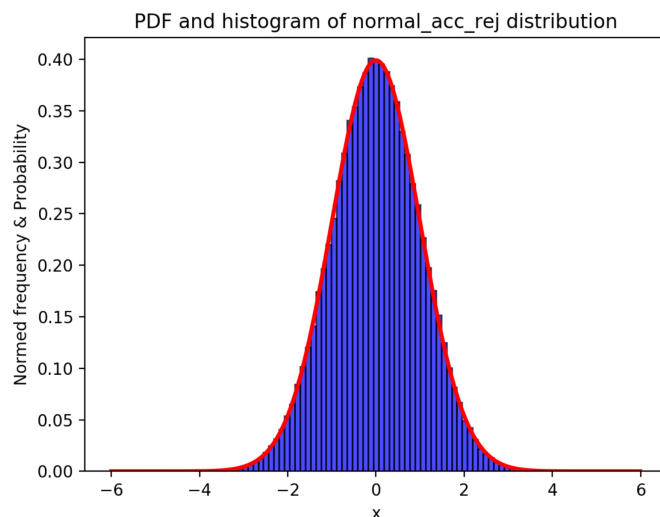


Figure 4: Normal distribution via Box-Muller method

3. Comparison

The results of 2 methods shown above are quite similar when $N = 200,000$. But the Box-Muller method may be more efficient, because Acceptance-Rejection method may reject many times when $u \geq \frac{p(y)}{cq(y)}$. Then, the recursive run of sampling code will take a lot of time.

4) Sampling from the Beta distribution (you can use any method introduced in our class). To sample from the Beta distribution from uniform distribution, I considered 2 different cases.

1. a, b greater than 1, and $a, b \in \mathbb{Z}^+$

We can refer to the relationship between beta distribution and order statistics. Notice that $u_j \sim \text{Beta}(j, n-j+1)$, then we can find the parameters n, j in order statistics can be represented by a, b as $j = a$ and $n = a + b - 1$. We can just create n i.i.d random variables which obey $U(0, 1)$, sort them and find j^{th} minimum of the sequence.

2. Otherwise

Notice that $Beta(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}$, where $\Gamma(x)$ is the standard gamma distribution $\Gamma(x, 1)$.

We can use Acceptance-Rejection method to get the result for gamma distribution, and combine them to sample from beta distribution.

The python code is shown below:

```
def beta(*args):
    assert(len(args) == 2)
    a, b = args[0], args[1]
    def order_stat(_n, _j):
        l = sorted([uniform(0,1) for _ in range(_n)])
        return l[_j-1]
    def beta_int(_a, _b):
        return order_stat(_a+_b-1, _a)
    if a < 1 and b < 1:
        u1, u2 = uniform(0, 1), uniform(0, 1)
        x = u1**(1/a); y = u2**(1/b)
        if x + y < 1:
            return x/(x+y)
        return beta(a, b)
    else:
        return beta_int(a, b)
```

For the PDFs, I use numeric integral to calculate $\beta(a, b)$:

```
def beta_integral(a, b):
    I, dx = 0, 0.001
    for i in np.arange(0.001, 1, dx):
        I += (i**(a-1))*((1-i)**(b-1))*dx
    return I
```

The function is tested under different cases: $Beta(0.5, 0.5)$, $Beta(2, 8)$, $Beta(1, 1)$ and $Beta(5, 5)$, all of these results are shown in Fig.5.

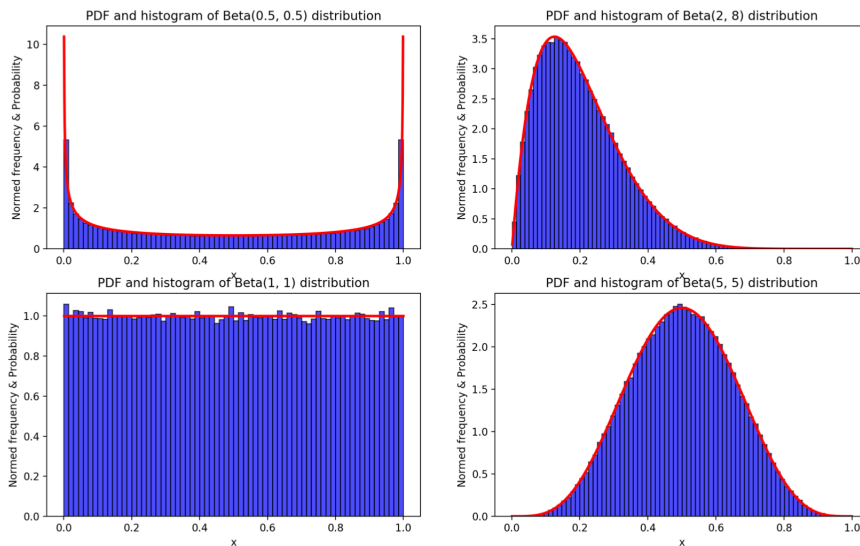


Figure 5: Beta distribution under different parameters

Problem 2

Given a random variable $X \sim N(0, 1)$, evaluate the tail probability $P(X > 8)$

- Use the standard sample average method.
- Use the importance sampling method.

Solution

- Standard Sample Average Method** Using standard sample average method, I generate 500,000 samples, count the number of items which is greater than 8. (The normal distribution samples are generated by Box-Muller method)

```
def std_sample():
    N, gt8 = 500000, 0
    samples = np.array([normal() for _ in range(N)])
    for i in range(N):
        if samples[i] > 8:
            gt8 += 1
    print("The tail probability evaluated by Standard Sample Average method \
        is: %.6f" % (gt8/N))
```

Unluckily, the tail probability is always 0 in many times of running the code shown above. The origin output from terminal is:

```
The tail probability evaluated by Standard Sample Average method is: 0.000000
```

- Importance Sampling Method** Using importance sampling method, let $\alpha = \mu = 8$, generate N samples x_k from $\mathcal{N}(\mu, 1)$, then by the importance sampling method, the tail probability can be evaluated as $P = \frac{1}{N} \sum_{x_k} I(x_k > 8) e^{\frac{1}{2}\mu^2 - \mu x_k}$. The python code is shown below:

```
def importance_sample():
    mu, N = 8, 500000
    samples = np.random.normal(mu, 1, N)
    P = sum(exp(0.5*mu**2 - mu*xk) for xk in samples if xk > mu) / N
    print("The tail probability evaluated by Importance sampling method \
        is: {}".format(P))
```

The output of runing code shown above is: $6.25247083608e-16$ The origin output from terminal is:

```
The tail probability evaluated by Importance sampling method is: 6.25247083608e-16
```

Problem 3

A coin with probability p of landing Heads is flipped repeatedly. Let N denote the number of flips until the pattern HH is observed.

- Suppose that p is a known constant, with $0 < p < 1$. Find $E(N)$
- Now suppose that p is unknown, and that we use a $Beta(a, b)$ prior to reflect our uncertainty about p (where a and b are known constants and are greater than 2). What is the expected number of flips until the pattern HH is observed.

Solution

- We can suppose the expectation of this event is $E(N)$.

- case 1. If the coin first landing with H, with $P(H) = p$.
 - case 1.1. The second landing of this coin is still H, with $P(HH) = p^2$
 - case 1.2. The second landing of this coin is T, with $P(HT) = p(1 - p)$. Now, we should restart this game.
- case 2. If the coin first landing with T, with $P(T) = 1 - p$. Under this case, we should restart this game.

Based on these analysis, we can get the following equation:

$$E(N) = 2p^2 + p(1 - p)(2 + E(N)) + (1 - p)(1 + E(N))$$

Hence,

$$E(N) = \frac{p + 1}{p^2}$$

b) We can solve this by using the conditional expectation and Adam's law:

$$E(N) = E(E(N|p)) = E\left(\frac{p + 1}{p^2}\right)$$

Then we need to find $E\left(\frac{p + 1}{p^2}\right)$.

First, find $E\left(\frac{1}{p}\right)$

$$\begin{aligned} E\left(\frac{1}{p}\right) &= \frac{1}{\beta(a, b)} \int_0^1 p^{a-2} (1 - p)^{b-1} dp \\ &= \frac{\beta(a - 1, b)}{\beta(a, b)} \\ &= \frac{a + b - 1}{a - 1} \end{aligned}$$

Then, calculate $E\left(\frac{1}{p^2}\right)$

$$\begin{aligned} E\left(\frac{1}{p^2}\right) &= \frac{1}{\beta(a, b)} \int_0^1 p^{a-3} (1 - p)^{b-1} dp \\ &= \frac{\beta(a - 2, b)}{\beta(a, b)} \\ &= \frac{(a + b - 1)(a + b - 2)}{(a - 1)(a - 2)} \end{aligned}$$

Hence, we can find the estimated expectation

$$E(N) = \frac{a + b - 1}{a - 1} + \frac{(a + b - 1)(a + b - 2)}{(a - 1)(a - 2)}$$

which equals to

$$E(N) = \frac{(a + b - 1)(2a + b - 4)}{(a - 2)(a - 1)}$$

Problem 4

Instead of predicting a single value for the parameter, we given an interval that is likely to contain the parameter: A $1 - \delta$ confidence interval for a parameter p is an interval $p \in [\hat{p} - \epsilon, \hat{p} + \epsilon]$ such that $Pr(p \in [\hat{p} - \epsilon, \hat{p} + \epsilon]) \geq 1 - \delta$. Now we toss a coin with probability p landing heads and probability $1 - p$ landing tails. The parameter p is unknown and we need to estimate its value from experiments results. We toss such coin N times, Let $X_i = 1$ if the i th result is head, otherwise 0. We estimate p by using $\hat{p} = \frac{X_1 + \dots + X_N}{N}$. Find the confidence interval for p , then discuss the impacts of δ and N .

Solution

First, by unbiased estimation:

$$E(\hat{p}) = \frac{1}{N} E\left(\sum_{i=1}^N X_i\right) = \frac{1}{N} Np = p$$

Then, for $0 \leq X_i \leq 1$, set $a = 0, b = 1$, then by Hoffding bound,

$$Pr(|\hat{p} - p| \geq \epsilon) = Pr\left(\left|\sum_{i=1}^N X_i - p\right| \geq \epsilon\right) \leq 2e^{-2N\epsilon^2}$$

Set $\delta = 2e^{-2N\epsilon^2}$, then

$$\epsilon = \sqrt{\frac{\ln(\frac{\delta}{2})}{2N}}$$

For $Pr(|\hat{p} - p| \geq \epsilon) \leq \delta$ is equals to $Pr(|\hat{p} - p| < \epsilon) > 1 - \delta$, which is the confidence interval

$$Pr(p \in (\hat{p} - \epsilon, \hat{p} + \epsilon)) > 1 - \delta$$

Hence,

$$\forall \delta > 0, \text{ with probability } 1 - \delta, |\hat{p} - p| < \sqrt{\frac{\ln(\frac{\delta}{2})}{2N}}$$

Impacts of δ and N : when N becomes larger, δ will also become larger. When N becomes larger, the length of the confidence interval will be smaller, and when δ becomes larger, the length of the confidence interval will be longer.

Problem 5

We know that the MMSE of X given Y is given by $g(Y) = E[X|Y]$. We also know that the Linear Least Square Estimate (LLSE) of X given Y , denoted by $L[X|Y]$, is shown as follows:

$$L(Y|X) = E(Y) + \frac{\text{cov}(X, Y)}{\text{Var}(Y)}(X - E(X))$$

Now we wish to estimate the probability of landing heads, denoted by θ , of a biased coin. We model θ as the value of a random variable θ with a known prior PDF $f_\Theta \sim \text{unif}(0, 1)$. We consider n independent tosses and let X be the number of heads observed. Find the MMSE $E[\Theta|X]$ and the LLSE $L[\Theta|X]$

Solution

- a) MMSE.

Because we set the prior PDF $p_\Theta \sim \text{unif}(0, 1) = \text{Beta}(1, 1)$, and when θ is given, k heads in n tosses is binomial distributed. So, by Beta-Binomial conjugacy,

$$\Theta|(x = k) \sim \text{Beta}(k + 1, n - k + 1)$$

Then, we can get the conditional expectation:

$$E(\Theta|x = k) = \frac{k+1}{n+2}$$

Hence, the MMSE is

$$E(\Theta|X) = \frac{X+1}{n+2}$$

- b) LLSE.

Because we set $\Theta \sim \text{unif}(0, 1)$, we can get:

$$\begin{aligned} E(\Theta) &= \frac{1}{2} \\ \text{Var}(\Theta) &= \frac{1}{12} \\ E(\Theta^2) &= \frac{1}{3} \end{aligned}$$

Also, when Θ is given to θ , $X|(\Theta = \theta)$ is a binomial distribution. Hence we can get:

$$\begin{aligned} E(X) &= E(E(X|\Theta)) = E(n\Theta) = \frac{n}{2} \\ \text{Var}(X) &= E(\text{Var}(X|\Theta)) + \text{Var}(E(X|\Theta)) = E(n\Theta(1-\Theta)) + \text{Var}(n\Theta) = \frac{n(n+2)}{12} \\ \text{Cov}(X, Y) &= E(\Theta X) - E(\Theta)E(X) = E(\Theta E(X|\Theta)) - E(\Theta)E(X) = \frac{n}{12} \end{aligned}$$

By the LLSE equation, we can get the following solution:

$$\begin{aligned} L(\Theta|X) &= E(\Theta) + \frac{\text{cov}(X, \Theta)}{\text{Var}(\Theta)}(X - E(X)) \\ &= \frac{1}{2} + \frac{\frac{n}{12}}{\frac{n(n+2)}{12}}(X - \frac{n}{2}) \\ &= \frac{X+1}{n+2} \end{aligned}$$

Problem 6

Given a coin with the probability p of landing heads. p is unknown and we need to estimate its value through data. In our data collection model, we have n independent tosses, result of each toss is either Head or Tail. Let X denote the number of heads in the total n tosses. Now we conduct experiments to collect data and find $X = k$. Then we need to find \hat{p} , the estimation of p .

- (a) Assume p is a random variable with a prior distribution $p \sim \text{Beta}(a, b)$, where a and b are known constants. Find \hat{p} through the MAP (Maximum a Posterior Probability) rule.
- (b) Assume p is an unknown constant. Find \hat{p} through the MLE (Maximum Likelihood Estimation) rule.
- (c) Assume p is a random variable with a prior distribution $p \sim \text{Beta}(a, b)$, where a and b are known constants. Find \hat{p} through the MMSE (Minimal Mean Squared Error) rule.

Solution

- a) MAP.

Recall the formula of MAP estimation:

$$\hat{\theta} = \arg \max_{\theta} p(\theta|X)$$

Because of $p(\theta|X) = \frac{p(X|\theta)p(\theta)}{p(X)}$, and $p(X)$ is independent with θ , Thus,

$$\hat{\theta} = \arg \max_{\theta} p(X|\theta)p(\theta)$$

which can be calculated as:

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} p(X|\theta)p(\theta) \\ &= \arg \max_{\theta} \theta^k (1-\theta)^{n-k} \cdot \frac{1}{\beta(a,b)} \theta^{a-1} (1-\theta)^{b-1}\end{aligned}$$

Take the partial derivatives of θ to 0,

$$\frac{\partial(\theta^k (1-\theta)^{n-k} \cdot \frac{1}{\beta(a,b)} \theta^{a-1} (1-\theta)^{b-1})}{\partial \theta} = 0$$

Hence,

$$\hat{\theta}_{MAP} = \frac{a+k-1}{a+b+n-2}$$

- b) MLE.

Let E be the event, we want to maximize $p(E|\theta)$, where

$$p(E|\theta) = \theta^k (1-\theta)^{n-k}$$

which equals to optimize the following equation

$$\begin{aligned}\hat{\theta}_{MLE} &= \arg \max_{\theta} \ln(p(E|\theta)) \\ &= \arg \max_{\theta} \ln(\theta^k (1-\theta)^{n-k})\end{aligned}$$

Take the derivative to 0,

$$\begin{aligned}&\frac{\partial (\ln(\theta^k (1-\theta)^{n-k}))}{\partial \theta} \\ &= \frac{\partial}{\partial \theta} (k \ln(\theta) + (n-k) \ln(1-\theta)) \\ &= \frac{k}{\theta} + \frac{n-k}{1-\theta} \\ &= 0\end{aligned}$$

Hence,

$$\hat{\theta}_{MLE} = \frac{k}{n}$$

- c) MMSE.

Because we set the prior PDF $p_{\Theta} \sim \text{Beta}(a, b)$, and when θ is given, k heads in n tosses is binomial distributed. So, by Beta-Binomial conjugacy,

$$\Theta|(x=k) \sim \text{Beta}(k+a, n-k+b)$$

Then, we can get the conditional expectation:

$$E(\Theta|x=k) = \frac{k+a}{n+a+b}$$

Hence, the MMSE is

$$\hat{\theta}_{MMSE} = E(\Theta|X) = \frac{k+a}{n+a+b}$$

Problem 7

Assume a random person's birthday is uniformly distributed on the 365 days of the year. People enter the room one by one. How many people are in the room the first time that two people share the same birthday? Let K be the desired number. Find $E(K)$ (integral form).

Solution

Consider there is only one person (i.e. $K = 1$), this probability is obviously zero. When the second person enter the room, the probability that they share the same birthday is $\frac{1}{365}$. So, the probability that $K = 2$ is $\frac{1}{365}$.

In more general cases, if $K = k$, which indicates the previous $k - 1$ people all have different birthdays. This event occurs with probability:

$$\begin{aligned} P_1 &= 1 \cdot \frac{364}{365} \cdot \frac{363}{365} \cdots \frac{365 - (k - 2)}{365} \\ &= \frac{365!}{365^{k-1}(365 - k)!} \end{aligned}$$

Then, the k^{th} person should share birthday with one of those $k - 1$ people, which has probability $P_2 = \frac{k-1}{365}$. So, the probability of $K = k$ is

$$P(K = k) = P_1 \cdot P_2 = \frac{k-1}{365} \cdot \frac{365!}{365^{k-1}(365 - k)!}$$

So, the expectation of K is

$$E(K) = \sum_{k=2}^{365} \frac{k-1}{365} \cdot \frac{365!}{365^{k-1}(365 - k)!} \cdot k$$

, the start $k = 2$ is for $p(K = 1) = 0$. (This result is the integral form). Moreover, to caculate the result, I write a python script to do some numeric calculations. The code is shown below:

```
def birthday_exp():
    E = 0
    for k in range(2, 365):
        prod = (k-1)/365
        for j in range(k-1):
            prod *= (365-j)/365
        E += k*prod
    return E
```

The expection calculated by this function is 24.616585894598874 (person).

Problem 8

Suppose buses arrive at a bus stop according to a Poisson process N_t with parameter λ . Given a fixed $t > 0$. The time of the last bus before t is S_{N_t} , and the time of the next bus after t is S_{N_t+1} . Show the following identity:

$$E(S_{N_t+1} - S_{N_t}) = \frac{2 - e^{-\lambda t}}{\lambda}$$

Solution

First, we compute $E(S_{N_t+1})$, by using conditional expectation $E(S_{N_t+1}|N_t = k)$,

$$E(S_{N_t+1}|N_t = k) = E(S_{k+1}|N_t = k)$$

, where S_{k+1} is independent with N_t , so,

$$E(S_{k+1}|N_t = k) = E(S_{k+1}) = \frac{k+1}{\lambda}$$

hence $E(S_{N_t+1}|N_t = k) = \frac{N_t+1}{\lambda}$. Taking expectation on both sides, and by Adam's law, we can get:

$$E(S_{N_t+1}) = E\left(\frac{N_t+1}{\lambda}\right) = t + \frac{1}{\lambda}$$

Then compute $E(S_{N_t})$, by conditional expectation $E(S_{N_t+1}|N_t = k)$. By the k^{th} arrival time is identical to k^{th} order statistics $u_{(k)}$, We can get:

$$\begin{aligned} E(S_k|N_t = k) &= E(u_{(k)}) \\ &= \int_0^\infty P(u_{(k)} > s) ds \\ &= \int_0^t \left(1 - \frac{s^k}{t^k}\right) ds \\ &= \frac{tk}{k+1} \end{aligned}$$

Hence,

$$E(S_{N_t}|N_t = k) = \frac{tN_t}{N_t+1}$$

Taking expectation on both sides, and by Adam's law, we can get:

$$\begin{aligned} E(S_{N_t}|N_t = k) &= \frac{tN_t}{N_t+1} \\ &= t - tE\left(\frac{1}{N_t+1}\right) \\ &= t - t \cdot \sum_{k=0}^{\infty} \frac{P(N_t = k)}{k+1} \\ &= t - t \cdot \sum_{k=0}^{\infty} \frac{e^{-\lambda t} (\lambda t)^k}{(k+1)k!} \\ &= t - t \cdot \left(\frac{e^{-\lambda t}}{\lambda t} (e^{-\lambda t} - 1)\right) \\ &= t - \frac{1}{\lambda} + \frac{e^{-\lambda t}}{\lambda} \end{aligned}$$

Thus,

$$\begin{aligned} E(S_{N_t+1} - S_{N_t}) &= \left(t + \frac{1}{\lambda}\right) - \left(t - \frac{1}{\lambda} + \frac{e^{-\lambda t}}{\lambda}\right) \\ &= \frac{2 - e^{-\lambda t}}{\lambda} \end{aligned}$$

Problem 9

Given k skill levels, we define a reward function $H(\cdot) = \{1, \dots, k\} \rightarrow R$. Then for skill levels $x \in \{1, \dots, k\}$ and $y \in \{1, \dots, k\}$, we define a soft-max function

$$\pi(x) = \frac{e^{H(x)}}{\sum_{y=1}^k e^{H(y)}}$$

Please show the following result:

for any skill level $x \in \{1, \dots, k\}$, we have

$$\frac{\partial \pi(x)}{\partial H(a)} = \pi(x)(1_{\{x=a\}} - \pi(a))$$

where 1_A is an index function of events, being 1 when event A is true and being 0 otherwise.

Solution

- a) If $x \neq a$:

$$\begin{aligned} \frac{\partial \pi(x)}{\partial H(a)} &= \frac{\partial \frac{e^{H(x)}}{\sum_{y=1}^k e^{H(y)}}}{\partial H(a)} \\ &= \frac{0 - e^{H(a)} e^{H(x)}}{(\sum_{y=1}^k e^{H(y)})^2} \\ &= -\frac{e^{H(a)}}{\sum_{y=1}^k e^{H(y)}} \cdot \frac{e^{H(x)}}{\sum_{y=1}^k e^{H(y)}} \\ &= -\pi(a)\pi(x) \\ &= \pi(x)(0 - \pi(a)) \end{aligned}$$

- b) If $x = a$:

$$\begin{aligned} \frac{\partial \pi(x)}{\partial H(a)} &= \frac{\partial \frac{e^{H(a)}}{\sum_{y=1}^k e^{H(y)}}}{\partial H(a)} \\ &= \frac{e^{H(a)} \cdot \sum_{y=1}^k e^{H(y)} - (e^{H(a)})^2}{(\sum_{y=1}^k e^{H(y)})^2} \\ &= \frac{e^{H(a)}}{\sum_{y=1}^k e^{H(y)}} \cdot \frac{\sum_{y=1}^k (e^{H(y)}) - e^{H(a)}}{\sum_{y=1}^k e^{H(y)}} \\ &= \pi(a)(1 - \pi(a)) \\ &= \pi(x)(1 - \pi(a)) \end{aligned}$$

Hence,

$$\begin{aligned} \frac{\partial \pi(x)}{\partial H(a)} &= \begin{cases} \pi(x)(0 - \pi(a)), & x \neq a \\ \pi(x)(1 - \pi(a)), & x = a \end{cases} \\ &= \pi(x)(1_{\{x=a\}} - \pi(a)) \end{aligned}$$