

Reinforcement Learning: Homework #2

Due on April 7, 2020 at 11:59pm

Professor Ziyu Shao

Tianyuan Wu
63305667

Problem 1

With the same format as bandit algorithms 1,2 and 3, write the pseudo-code of gradient bandit algorithm for this three-armed Bernoulli bandit problem.

Solution

Algorithm 1 Gradient Algorithm

Initialize $\bar{R}_t = 0, H(j) = 0, j \in \{1, 2, 3\}$

```

1: for  $t = 1, 2, 3, \dots, N$  do
2:    $\pi_t(j) \leftarrow \frac{e^{H(j)}}{\sum_{k=1,2,3} e^{H(k)}}, j \in \{1, 2, 3\}$ 
3:   Sample  $I(t)$  with probability  $P(I(t) = j) = \pi_t(j)$ 
4:    $\bar{R}_t \leftarrow \frac{1}{t} r_{I(t)} + \frac{t-1}{t} \bar{R}_t$ 
5:   if Baseline is set then
6:      $B_t \leftarrow \text{Baseline}$ 
7:   else
8:      $B_t \leftarrow \bar{R}_t$ 
9:   end if
10:   $H(I(t)) \leftarrow H(I(t)) + \alpha(r_{I(t)} - B_t)(1 - \pi_t(I(t)))$ 
11:   $H(i) \leftarrow H(i) - \alpha(r_{I(t)} - B_t)\pi_t(i)$ , Forall  $i \in \{1, 2, 3\}$  and  $i \neq I(t)$ 
12: end for
```

Problem 2

Now suppose we obtain the Bernoulli distribution parameters from an oracle, which are shown in the following table below. Choose $N = 10000$ and compute the theoretically maximized expectation of aggregate rewards over N time slots. We call it the oracle value. Note that these parameters $\theta_j, j = 1, 2, 3$ and oracle values are unknown to all bandit algorithms.

Arm_j	1	2	3
θ_j	0.9	0.8	0.7

Solution

We assume Arm_1 is chosen x times, Arm_2 is chosen y times, then Arm_3 is chosen $10000 - x - y$ times. By the expectation of binomial distribution:

$$E(X) = np, \quad X \sim \text{Bin}(n, p)$$

We can obtain the expectation of aggregate reward \mathbb{E}_{aggr}

$$\mathbb{E}_{aggr} = 0.9x + 0.8y + 0.7(10000 - x - y), \quad s.t. \ x, y \in \mathbb{N}, x + y \leq 10000$$

Hence, take $x = 10000$ and $y = 0$, we can get

$$\begin{aligned}
 \mathbb{E}_{aggr} &= 0.9x + 0.7(10000 - x) \\
 &\leq 0.9 \cdot 10000 + 0.7 \cdot 0 \\
 &= 9000
 \end{aligned}$$

Thus, if $N = 10000$, the theoretically maximized expectation is 9000.

Problem 3

Implement classical bandit algorithms

Solution

Codes are in the jupyter-notebook document (.ipynb file).

Problem 4

Each experiment lasts for $N = 5000$ turns, and we run each experiment 1000 times. Results are averaged over these 1000 independent runs.

Solution

Algorithm	Parameters	Results			
		$\hat{\theta}_1$	$\hat{\theta}_2$	$\hat{\theta}_3$	Aggregate Reward
ϵ -greedy	$\epsilon = 0.1$	0.9002	0.8008	0.7010	4446.6
ϵ -greedy	$\epsilon = 0.5$	0.9002	0.7995	0.7000	4248.7
ϵ -greedy	$\epsilon = 0.9$	0.8995	0.8005	0.6990	4047.7
UCB	$c = 1$	0.9001	0.7966	0.6921	4386.7
UCB	$c = 5$	0.9000	0.7997	0.6991	4124.6
UCB	$c = 10$	0.8999	0.7998	0.6999	4062.8
TS	$\{1, 1\}, \{1, 1\}, \{1, 1\}$	0.8998	0.7519	0.6282	4485.1
TS	$\{601, 401\}, \{401, 601\}, \{2, 3\}$	0.6835	0.4002	0.6703	3776.2
Gradient	baseline=0	-	-	-	4425.2
Gradient	baseline=0.8	-	-	-	4431.1
Gradient	baseline=5	-	-	-	4200.5
Gradient	baseline=20	-	-	-	4109.6
Gradient	$\beta = 0.2$	-	-	-	4256.6
Gradient	$\beta = 1$	-	-	-	4428.4
Gradient	$\beta = 2$	-	-	-	4460.1
Gradient	$\beta = 5$	-	-	-	4481.0
Gradient	$\beta = \frac{5(N-t)}{t}$	-	-	-	4490.1

For the time-variant algorithm, I take $\beta = \frac{2(N-t)}{t}$. And for all gradient algorithms, I set $\alpha = 0.1$.

Problem 5

Solution

1. ϵ -greedy Algorithm

Overall Regret

Table 1: ϵ -greedy Algorithm

Parameter	Accumulated Avg. Regret	Total percentage of optimals
$\epsilon=0.9$	449.916	39.94%
$\epsilon=0.5$	250.978	66.28%
$\epsilon=0.1$	53.508	92.52%

Regret-time function

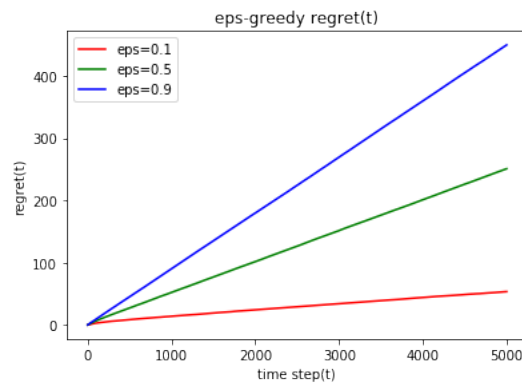


Figure 1: ϵ -greedy algorithm: regret-t

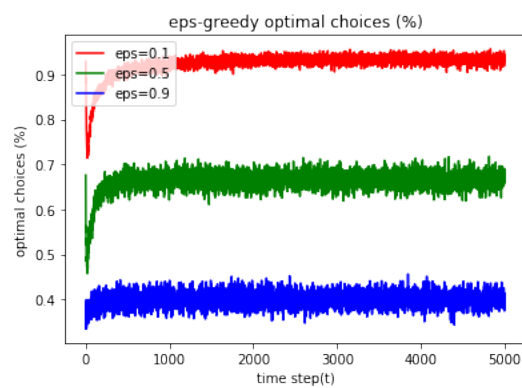


Figure 2: ϵ -greedy algorithm: Percentage of optimals

2. UCB Algorithm

Overall Regret

Table 2: UCB Algorithm

Parameter	Accumulated Avg. Regret	Total percentage of optimals
c=1	114.09	82.44%
c=5	374.568	47.28%
c=10	437.854	40.20%


Regret-time function

Figure 3: UCB algorithm: regret-t

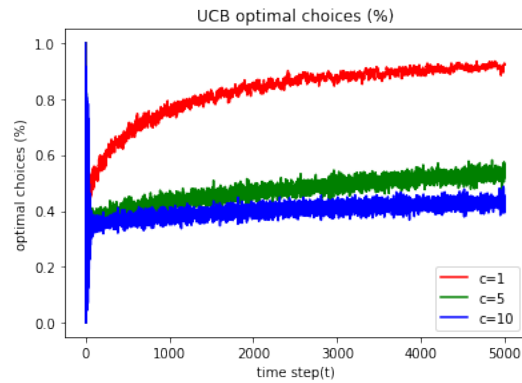


Figure 4: UCB algorithm: Percentage of optimals

3. TS Algorithm

Overall Regret

Table 3: TS Algorithm

Parameter	Accumulated Avg. Regret	Total percentage of optimals
$\{\{1,1\}, \{1,1\}, \{1,1\}\}$	15.985	97.55%
$\{\{601,401\}, \{401,601\}, \{2,3\}\}$	723.78	28.67%


Regret-time function

Figure 5: TS algorithm: regret-t

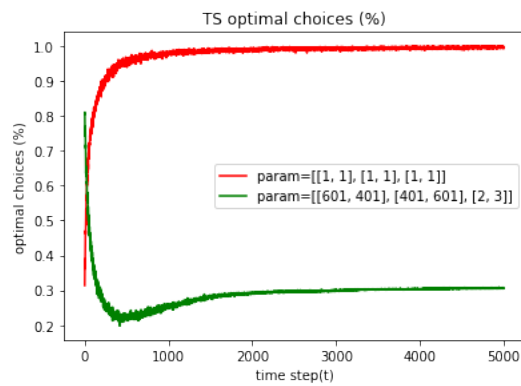


Figure 6: TS algorithm: Percentage of optimals

4. Gradient algorithm

Overall Regret

Table 4: Gradient Algorithm

Parameter	Accumulated Avg. Regret	Total percentage of optimals
baseline=0	74.8	89.83%
baseline=0.8	68.9	90.66%
baseline=5	299.5	60.14%
baseline=20	390.4	47.98%
$\beta = 0.2$	243.4	71.64%
$\beta = 1$	71.6	90.62%
$\beta = 2$	39.9	94.66%
$\beta = 5$	19.0	97.43%
time-variant	9.899	98.68%


Regret-time function

Figure 7: Gradient algorithm (baseline): regret-t




Figure 8: Gradient algorithm (baseline): Percentage of optimals




Figure 9: Gradient algorithm (parameter): regret-t




Figure 10: Gradient algorithm (parameter): Percentage of optimals




Figure 11: Gradient algorithm (time-varient): regret-t




Figure 12: Gradient algorithm (time-varient): Percentage of optimals

Problem 6

Algorithm	Parameter	Gap
ϵ -greedy	$\epsilon = 0.1$	53.5
ϵ -greedy	$\epsilon = 0.5$	251.0
ϵ -greedy	$\epsilon = 0.9$	449.9
UCB	$c = 1$	114.1
UCB	$c = 5$	374.6
UCB	$c = 10$	437.8
TS	$\{1, 1\}, \{1, 1\}, \{1, 1\}$	15.9
TS	$\{601, 401\}, \{401, 601\}, \{2, 3\}$	723.8
Gradient	baseline=0	74.8
Gradient	baseline=0.8	68.9
Gradient	baseline=5	299.5
Gradient	baseline=20	390.4
Gradient	$\beta = 0.2$	243.4
Gradient	$\beta = 1$	71.6
Gradient	$\beta = 2$	39.9
Gradient	$\beta = 5$	19.0
Gradient	time-varient	9.9

The results are shown in Table.x, we can observe that the gradient algorithm with time-varient β has the best performance. The impacts are:

1. ϵ -greedy

When ϵ becomes larger, the exploitation increases, while exploration decreases.

When ϵ becomes smaller, the exploitation decreases, while exploration increases.

So, during the process that ϵ increases from a small value to 1, the gap first becomes smaller, then becomes larger.

2. UCB

When c increases, the gap becomes larger.

3. TS

When the parameters satisfy $\frac{\alpha_i}{\alpha_i + \beta_i} = \theta_i$, the gap reaches the smallest value.

4. Gradient

When baseline reaches the optimal value (0.9 in this situation), the gap reaches the smallest value.

Also, during the process that β increases from a small value to a large value, the gap first becomes smaller, then becomes larger.

Problem 7

Give your understanding of the exploration-exploitation trade-off in bandit algorithms.

Solution

Exploration is to “try something new”, is to explore if there are something better to choose, it may have better rewards, also may have worse rewards. In bandit algorithm, it's to choose a new bandit, it may different from the current optimal one. The algorithm doesn't know it will be better or worse. For example, in ϵ -greedy algorithm, the ϵ probability is to explore, prevent the algorithm from locking on the local optima. Exploitation is to choose the current optimal bandit based on observations/experiments. It's a greedy choice. So, if we only prefer exploitation, the algorithm will make choice randomly, cannot converge to the optimal value. But if we only consider about exploration, the algorithm may “lock” on some local optimals, cannot find the global optimal solution. Thus, we need to balance exploration and exploitation, try to explore more in the beginning, and then make greedy choices to converge to the optimal value.

Problem 8

We implicitly assume the reward distribution of three arms are independent. How about the dependent case? Can you design an algorithm to exploit such information to obtain a better result?

Solution

We can refer to TLP(two-level policy) algorithm, which is introduced in “multi-armed bandit problems with dependent arms - Sandeep Pandey, Deepayan Chakrebari and Deepak Agarwal, Yahoo! research”. For example, we have 3 arms and a prior that $\theta_1 - \theta_2 \leq 0.01$, then we can divide arm 1,2 into cluster1, and arm3 to cluster2. Thus, we can construct a model: n arms are divided into K clusters, by exploiting the dependence information. Then, for each step, we calculate the average reward \bar{r}_k and variance $\sigma(r_k)$ for each cluster. Then call the independent bandit algorithm (e.g. UCB, gradient, etc.) to select a cluster. Then, in the selected group, call independent bandit algorithm again to get a choice. It can be described in pseudo-code:

Algorithm 2 TLP Algorithm

- 1: Divide Arm_1, \dots, Arm_n into K clusters C_1, \dots, C_K by dependence knowledge.
 - 2: **for** $t = 1, 2, 3, \dots, N$ **do**
 - 3: Foreach cluster C_i , calculate $\hat{\sigma}_{C_i}(t)$, $\hat{r}_{C_i}(t)$.
 - 4: Call independent bandit algorithm to choose a cluster C_t
 - 5: Call independent bandit algorithm in C_t , to choose Arm_t
 - 6: **end for**
-

Problem 9

We can simplify this problem by sum up the cost of choosing some arm and the reward of choosing it to a overall reward. That is, let

$$R_{overall, arm_i} = R_{arm_i} - cost_{arm_i}$$

Then, we can use any bandit algorithm without constraint, and maximize the overall reward. For example, if we use epsilon-greedy algorithm, then the pseudo-code is shown below:

Algorithm 3 ϵ -greedy algorithm for constraint bandit

```

1: Let  $Arm_1, \dots, Arm_n$  be the  $n$  arms
2:  $overall = \{0, 0, \dots, 0\}$ 
3: for  $t = 1, 2, 3, \dots, N$  do
4:   if  $unif(0, 1) < \epsilon$  then
5:      $arm = \text{random\_choice}(\text{arms})$ 
6:   else
7:      $arm = \text{argmax}(\text{arms}, \text{overall})$ 
8:   end if
9:    $\text{reward}, \text{cost} = \text{bandit}[arm].\text{play}()$ 
10:   $\text{overall}[arm] = \text{reward} - \text{cost}$ 
11: end for

```

Problem 10

Please reproduce the proof of regret decomposition lemma.

Solution

For $Q(a_\tau) = \mathbb{E}[r_\tau | a_\tau]$ is a random variable, and by adam's rule,

$$\mathbb{E}[Q(a_\tau)] = \mathbb{E}[\mathbb{E}[r_\tau | a_\tau]] = \mathbb{E}[r_\tau]$$

Let $S_t = \sum_{\tau=1}^t Q(a_\tau)$, then we have:

$$\mathbb{E}(S_t) = \sum_{\tau=1}^t \mathbb{E}[Q(a_\tau)] = \mathbb{E}\left[\sum_{\tau=1}^t r_\tau\right]$$

Since $\sum_{a \in A} 1_{a_\tau=a} = 1$ for any fixed τ , then we have:

$$\begin{aligned}
\mathbb{E}(S_t) &= \mathbb{E}\left[\sum_{\tau=1}^t r_\tau\right] \\
&= \mathbb{E}\left[\sum_{\tau=1}^t \sum_{a \in A} r_\tau 1_{a_\tau=a}\right] \\
&= \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[r_\tau 1_{a_\tau=a}]
\end{aligned}$$

On the other hand, $\sum_{\tau=1}^t \sum_{a \in A} 1_{a_\tau=a} = t$, so $\sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[1_{a_\tau=a}] = t$.
The total regret can be calculated as:

$$\begin{aligned}
L_t &= \mathbb{E}[\sum_{\tau=1}^t V^* - Q(a_\tau)] \\
&= tv^* - \mathbb{E}(\sum_{\tau=1}^t Q(a_\tau)) \\
&= tv^* - \mathbb{E}(S_t) \\
&= tv^* - \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[r_\tau 1_{a_\tau=a}] \\
&= \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[1_{a_\tau=a}] v^* - \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[r_\tau 1_{a_\tau=a}] \\
&= \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[(v^* - r_\tau) 1_{a_\tau=a}]
\end{aligned}$$

Then, we calculate $\mathbb{E}[(v^* - r_\tau) 1_{a_\tau=a} | a_\tau]$:

$$\begin{aligned}
&\mathbb{E}[(v^* - r_\tau) 1_{a_\tau=a} | a_\tau] \\
&= 1_{a_\tau=a} \mathbb{E}[(v^* - r_\tau) | a_\tau] \\
&= 1_{a_\tau=a} (v^* - Q(a)) \\
&= 1_{a_\tau=a} \Delta_a
\end{aligned}$$

Thus,

$$\mathbb{E}[(v^* - r_\tau) 1_{a_\tau=a}] = \mathbb{E}[1_{a_\tau=a} \Delta_a]$$

Then we have:

$$\begin{aligned}
L_t &= \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[(v^* - r_\tau) 1_{a_\tau=a}] \\
&= \sum_{a \in A} \sum_{\tau=1}^t \mathbb{E}[1_{a_\tau=a} \Delta_a] \\
&= \sum_{a \in A} \mathbb{E}[\sum_{\tau=1}^t 1_{a_\tau=a}] \Delta_a \\
&= \sum_{a \in A} \mathbb{E}[N_t(a)] \Delta_a
\end{aligned}$$

Problem 11

Please reproduce the derivation of gradient bandit algorithm.

Solution

For the gradient ascent algorithm, for each iteration, we have:

$$w_{t+1} \leftarrow w_t + \nabla_w \mathbb{E}[f(w_t)]$$

And the objective function is $\max \mathbb{E}[R_t]$. We have:

$$\max \mathbb{E}[R_t] = \max \mathbb{E}[\mathbb{E}[R_t | E_t]] = \sum_x q_*(x) \pi_t(x)$$

where $\pi_t(x) = \frac{\exp(H_t(x))}{\sum_{y=1}^k \exp(H_t(y))}$. So, $\mathbb{E}[R_t]$ is a function of $H_t(x)$. And for each iteration, we have:

$$H_{t+1}(a) = H_t(a) + \alpha \frac{\partial \mathbb{E}[R_t]}{\partial H_t(a)}$$

Then we prove the lemma:

$$\frac{\partial \pi(x)}{\partial H(a)} = \pi(x)(1_{\{x=a\}} - \pi(a))$$

Proof:

- a) If $x \neq a$:

$$\begin{aligned} \frac{\partial \pi(x)}{\partial H(a)} &= \frac{\partial \frac{e^{H(x)}}{\sum_{y=1}^k e^{H(y)}}}{\partial H(a)} \\ &= \frac{0 - e^{H(a)} e^{H(x)}}{(\sum_{y=1}^k e^{H(y)})^2} \\ &= -\frac{e^{H(a)}}{\sum_{y=1}^k e^{H(y)}} \cdot \frac{e^{H(x)}}{\sum_{y=1}^k e^{H(y)}} \\ &= -\pi(a)\pi(x) \\ &= \pi(x)(0 - \pi(a)) \end{aligned}$$

- b) If $x = a$:

$$\begin{aligned} \frac{\partial \pi(x)}{\partial H(a)} &= \frac{\partial \frac{e^{H(a)}}{\sum_{y=1}^k e^{H(y)}}}{\partial H(a)} \\ &= \frac{e^{H(a)} \cdot \sum_{y=1}^k e^{H(y)} - (e^{H(a)})^2}{(\sum_{y=1}^k e^{H(y)})^2} \\ &= \frac{e^{H(a)}}{\sum_{y=1}^k e^{H(y)}} \cdot \frac{\sum_{y=1}^k (e^{H(y)}) - e^{H(a)}}{\sum_{y=1}^k e^{H(y)}} \\ &= \pi(a)(1 - \pi(a)) \\ &= \pi(x)(1 - \pi(a)) \end{aligned}$$

Hence,

$$\begin{aligned} \frac{\partial \pi(x)}{\partial H(a)} &= \begin{cases} \pi(x)(0 - \pi(a)), & x \neq a \\ \pi(x)(1 - \pi(a)), & x = a \end{cases} \\ &= \pi(x)(1_{\{x=a\}} - \pi(a)) \end{aligned}$$

Hence,

$$\frac{\partial \pi_t(A_t)}{\partial H_t(a)} = \pi_t(x)(1_{\{A_t=a\}} - \pi_t(a))$$

and then:

$$\frac{\partial \mathbb{E}[R_t]}{\partial H_t(a)} = \sum_x q_*(x) \frac{\partial \pi_t(A_t)}{\partial H_t(a)}$$

Since $\sum_x \pi_t(x) = 1$, we have $\sum_x \frac{\partial \pi_t(A_t)}{\partial H_t(a)} = 0$.

Thus, we introduce a baseline B_t , which is not depend on x , satisfies $\sum_x B_t \frac{\partial \pi_t(A_t)}{\partial H_t(a)} = 0$.

Then we calculate the partial derivative $\frac{\partial \mathbb{E}(R_t)}{\partial H_t(a)}$

$$\begin{aligned}
\frac{\partial \mathbb{E}(R_t)}{\partial H_t(a)} &= \sum_x (q_*(x) - B_t) \frac{\partial \pi_t(A_t)}{\partial H_t(a)} \\
&= \sum_x \pi_t(q_*(x) - B_t) \frac{\partial \pi_t(A_t)}{\partial H_t(a)} \frac{1}{\pi_t} \\
&= \sum_x \pi_t(q_*(x) - B_t) (1_{x=a} - \pi_t(a)) \\
&= \mathbb{E}[(q_*(x) - B_t)(1_{x=a} - \pi_t(a))]
\end{aligned}$$

Now we choose $B_t = \bar{R}_t = \frac{1}{t} \sum_{i=1}^t R_i$.

Then we will show:

$$\mathbb{E}[q_*(A_t)(1_{x=a} - \pi_t(a))] = \mathbb{E}[R_t(1_{x=a} - \pi_t(a))]$$

Proof:

$$\begin{aligned}
&\mathbb{E}[(q_*(A_t))(1_{x=a} - \pi_t(a))] \\
&= \mathbb{E}[\mathbb{E}[R_t|A_t](1_{A_t=a} - \pi_t(a))] \\
&= \mathbb{E}[\mathbb{E}[R_t(1_{A_t=a} - \pi_t(a))|A_t]] \\
&= \mathbb{E}[R_t(1_{A_t=a} - \pi_t(a))]
\end{aligned}$$

Thus we have:

$$\frac{\partial \mathbb{E}(R_t)}{\partial H_t(a)} = \mathbb{E}[(q_*(A_t) - \bar{R}_t)(1_{A_t=a} - \pi_t(a))] = \mathbb{E}[(R_t - \bar{R}_t)(1_{A_t=a} - \pi_t(a))]$$

Then, we have the gradient bandit algorithm:

$$H_{t+1}(a) = H_t(a) + \alpha(R_t - \bar{R}_t)(1_{A_t=a} - \pi_t(a)), \quad \forall \alpha \in \{1, \dots, k\}$$

And for the policy gradient, we want to maximize $\mathbb{E}_x[f(x)]$, where $x \sim P_\theta$, P_θ is the policy, we have:

$$\begin{aligned}
\nabla_\theta \mathbb{E}[f(x)] &= \nabla_\theta \sum_x p(x) f(x) \\
&= \sum_x p(x) \frac{\nabla_\theta p(x)}{p(x)} f(x) \\
&= \sum_x p(x) (\nabla_\theta \log(p(x))) f(x) \\
&= \mathbb{E}[f(x) \nabla_\theta \log(p(x))]
\end{aligned}$$